

Semantic Chunking and Hybrid Retrieval for Financial Document QA: An Ablation Study on FinanceBench

Ebuka Chidubem Prudence Uzoama

Anote AI Research Fellowship 2026, Supervised by Natan Vidra, CEO, Anote AI

Abstract

Retrieval augmented generation pipelines for financial document question answering are commonly enhanced with semantic chunking, cross-encoder reranking, metadata enrichment, query expansion, and hybrid combinations of these, on the expectation that each one improves on a plain fixed-chunking baseline. This paper tests that expectation directly with a controlled three-condition ablation on a stratified 14-question sample of FinanceBench: a fixed-chunking baseline, semantic-boundary chunking, and a hybrid pipeline that combines structure-aware chunking, BM25 and TF-IDF fusion, metadata prefixes, HyDE-style query expansion, and lexical reranking. All three conditions use Claude as the generation model, and retrieval throughout is built on TF-IDF and BM25 rather than neural embeddings and a cross-encoder reranker. To guard against a single run being over-interpreted, the entire pipeline was independently repeated end to end, with a second, separately written set of answers scored the same way. Neither semantic chunking nor the hybrid pipeline beats the baseline in either run. The semantic-chunking shortfall is directionally consistent and sits right at the edge of statistical significance across the two runs, significant in one and just short of it in the other. The hybrid shortfall is smaller and not significant in either run. Exact Match could not be meaningfully evaluated at this sample size, since every metrics-generated question failed retrieval outright regardless of condition.

Keywords: retrieval-augmented generation, financial document QA, chunking strategies, FinanceBench, reproducibility, negative results.

1. Introduction

Retrieval augmented generation is the dominant approach for grounding large language model answers in financial filings, and practitioners commonly reach for enhancements such as semantic chunking, cross-encoder reranking, metadata enrichment, query expansion, and hybrid combinations of these, expecting each one to improve on a plain fixed-chunking baseline. Whether that expectation holds on financial documents specifically is an open, practically important question. SEC filings pack dense numeric tables, multi-level headers, and cross-references into a format that general-purpose RAG benchmarks do not stress in the same way (Setty et al., 2024; Akarsu et al., 2026).

This paper runs a controlled ablation that tests that expectation directly, comparing a fixed-chunking baseline against semantic-boundary chunking and a full hybrid retrieval pipeline on a stratified sample of FinanceBench questions. Answers throughout are generated by Claude, and retrieval is built on TF-IDF and BM25 rather than neural embeddings and a cross-encoder reranker, so the comparison between conditions is internally consistent even though the raw scores are specific to this particular model and retrieval stack.

Because a single run can make small differences look more solid than they are, this paper reports two independent runs of the entire pipeline: source documents downloaded again, retrieval rerun, and answers written a second time without reference to the first run's text.

This paper makes the following contributions:

- A controlled three-condition ablation, fixed-chunking baseline, semantic-boundary chunking, and a hybrid pipeline, on FinanceBench.

- Two independent end-to-end runs of the same pipeline, checking whether the findings hold up to the ordinary variance of free-text answer generation and to the retrieval pipeline's own reproducibility.
- A finding that neither enhancement beats the baseline in either run: the semantic-chunking shortfall sits right at the edge of statistical significance, and the hybrid shortfall is smaller and not significant.
- A full reproducibility trail, so every number in this paper can be traced back to what actually produced it.

2. Related Work

2.1 Financial Document Question Answering

FinanceBench (Islam et al., 2023) is the benchmark used throughout: 10,231 questions over SEC filings, of which a 150-question public subset is used here. Chen et al. (2021) introduce FinQA, establishing that multi-step numerical reasoning over financial tables is difficult even for strong models, consistent with the near-zero ROUGE-L scores on metrics-generated questions reported here.

2.2 RAG for Financial Documents

Setty et al. (2024) test chunking, query expansion, metadata, and reranking on FinanceBench with GPT-4-Turbo, and provide the closest published precedent for the ablation design used here. Choi et al. (2025) introduce FinAgentBench, which treats each financial table as a single retrievable unit precisely because standard chunking breaks table structure, a failure mode consistent with the near-zero scores on metrics-generated questions observed in every condition tested in this paper. Akarsu et al. (2026) benchmark ten retrieval strategies on mixed text-and-table financial documents and find that BM25 outperforms dense retrieval on this document type, directionally consistent with this paper's own BM25-based hybrid condition. Dadopoulos et al. (2025) test metadata-enriched chunking on FinanceBench and find that embedding metadata directly into chunk text helps more than post-retrieval reranking alone.

2.3 Embeddings and Metric Selection

Anderson et al. (2024) show that finance-specific text embeddings improve FinanceBench accuracy by roughly 8 percentage points over general-purpose embeddings. This paper's use of TF-IDF instead of neural embeddings should be read against that result. TF-IDF is weaker than even general-purpose embeddings, so the retrieval quality tested here is a conservative floor rather than a realistic estimate of what a stronger dense-retrieval system would achieve.

3. Methodology

3.1 Dataset

Questions are drawn from the publicly available FinanceBench open-source subset, 150 questions spanning metrics-generated, novel-generated, and domain-relevant question types, after excluding 17 questions whose source PDFs are not present in the public repository (Adobe 10-Ks 2015, 2016, 2017, and 2022; Johnson and Johnson's 2022 10-K and three 2022 to 2023 filings; MGM Resorts' Q4 2022 earnings). From the remaining pool, a stratified sample of 15 questions was drawn with a fixed seed, 5 per question type. One sampled filing, AMD's 2015 10-K, failed to parse under two different PDF text extractors and was dropped, leaving 14 questions across 11 filings: 5 domain-relevant, 5 novel-generated, and 4 metrics-generated.

3.2 Conditions Tested

Three conditions spanning a natural range of complexity are tested: a minimal fixed-chunking baseline, a single low-cost enhancement in semantic-boundary chunking, and a maximal combined enhancement in a hybrid pipeline. All three use the same generation model so the comparison between them is fair.

Condition	Chunking	Reranking	Metadata	Query Exp.
Baseline	Fixed, word-based	none	none	none
Semantic-boundary	Paragraph / sentence	none	none	none
Hybrid	Recursive, structured	Lexical overlap	Prefix	HyDE-style

Table 1: Conditions tested in this paper. All three use Claude as the generation model.

Pipeline Components

Baseline. Fixed-length chunking approximating 512-token, 50-token-overlap splitting, implemented as a 380-word, 40-word-overlap word-based split. The top 4 chunks by TF-IDF cosine similarity form the context.

Semantic-boundary chunking. Paragraph-boundary splitting, falling back to sentence boundaries where a filing lacks blank-line structure, with a 1,024-character cap per chunk.

Hybrid. Recursive, structure-aware chunking (2,000-character target, 200-character overlap) with a company, year, and filing-type metadata prefix on each chunk. BM25 and TF-IDF are combined by equal-weight reciprocal-rank fusion over a HyDE-style expanded query, then reordered by a lexical-overlap score and truncated to the top 5 chunks.

3.3 Reproducibility: Two Independent Runs

To check that the first run's finding was not an artifact of a stale or mis-saved intermediate file, the entire pipeline was rebuilt from scratch for an independent second pass. The FinanceBench open-source PDFs were downloaded again, the eleven filings needed for this paper's sample were re-extracted, and the retrieval pipeline was rerun unmodified against the freshly downloaded documents. The regenerated retrieval contexts matched the first run's contexts closely, in most cases exactly, since TF-IDF and BM25 are deterministic over fixed text. That consistency is itself informative: the pipeline is a genuine, reproducible procedure rather than a one-off result. A second, independent set of answers was then written fresh from the regenerated context, without reference to the first run's answer text, and rescored with the same procedure. Both runs are reported throughout this paper.

4. Evaluation Metrics

- ROUGE-L F1: with LCS(pred, gold) the longest common subsequence length between predicted and gold answer tokens, precision $P = \text{LCS} / |\text{pred}|$, recall $R = \text{LCS} / |\text{gold}|$, $F1 = 2PR / (P + R)$, computed with stemming.
- Exact Match (EM): 1 if the normalized gold answer appears as a substring of the normalized prediction, else 0.
- Bootstrap 95% CI: the per-question F1 difference between two conditions is resampled with replacement, 5,000 resamples in Run 1 and 10,000 in Run 2, and the 2.5th and 97.5th percentiles of the resampled means form the interval.

5. Results

5.1 Overall

Condition	Run 1 F1	Run 2 F1
Baseline	0.152	0.145
Semantic (change from baseline)	0.117 (-0.035)	0.123 (-0.023)
Hybrid (change from baseline)	0.128 (-0.024)	0.124 (-0.021)

Table 2: Mean ROUGE-L F1 by condition, across two independent end-to-end runs.

5.2 Results by Question Type (Run 1)

Question type	Baseline F1	Semantic F1	Hybrid F1	Baseline EM	Semantic EM	Hybrid EM
Domain-relevant	0.242	0.194	0.224	0.000	0.000	0.000
Metrics-generated	0.026	0.000	0.000	0.000	0.000	0.000
Novel-generated	0.163	0.133	0.133	0.200	0.000	0.200

Table 3: ROUGE-L F1 and Exact Match by question type and condition, Run 1.

Metrics-generated questions score at or near zero on both ROUGE-L and Exact Match across all three conditions. All four sampled metrics-generated questions failed retrieval outright in every condition (Appendix, Table A1), so there is no case in this sample where a correct numeric answer existed to be measured by either metric.

5.3 Difference from Baseline: Two Independent Bootstraps

Comparison	Run 1 diff	Run 1 95% CI	Run 2 diff	Run 2 95% CI
Semantic minus baseline	-0.035	(-0.072, -0.002)	-0.023	(-0.049, +0.001)
Hybrid minus baseline	-0.025	(-0.071, +0.019)	-0.021	(-0.066, +0.020)

Table 4: Paired bootstrap comparison against baseline, both independent runs.

The semantic-chunking shortfall crosses the significance threshold in Run 1 and sits just outside it in Run 2, with an upper bound of +0.0008, effectively zero. Taken together, this looks like a real but fragile effect at this sample size, sensitive to the ordinary wording variance between two independent generation passes. The hybrid condition is not significant in either run.

Finding 1: The semantic-chunking condition's shortfall relative to baseline shows up in both independent runs, with a similar size in each (-0.035 in Run 1, -0.023 in Run 2). Its statistical significance is fragile: significant in Run 1, just short of significant in Run 2.

Finding 2: The hybrid condition also scores below baseline in both independent runs (-0.025, -0.021), though the difference is not statistically significant in either, despite the hybrid combining every enhancement tested here at once.

Finding 3: Exact Match cannot be meaningfully evaluated at this sample size. Of the four metrics-generated questions sampled here, retrieval failed outright in every condition, so Exact Match is zero throughout and does not distinguish the conditions.

6. Discussion

Two patterns stand out. First, semantic-boundary chunking's shortfall relative to baseline is consistent across two independent generation passes, and similar in size in both, which suggests it reflects something about how paragraph- and sentence-boundary chunks interact with the retrieval method used here, rather than pure generation noise. Semantic chunks in this pipeline are longer and more topically diffuse than the fixed-length baseline chunks, which may make them harder for TF-IDF's exact-term matching to rank correctly against a short, specific query.

Second, the hybrid condition does not recover the baseline's performance in either run, even though it combines every enhancement tested here at once: structure-aware chunking, metadata prefixes, reciprocal-rank fusion of BM25 and TF-IDF, HyDE-style query expansion, and lexical reranking. Combining more retrieval techniques does not automatically compound into better retrieval. If each individual technique is already a simplified stand-in for a more capable tool, stacking several of them can add as much noise as it removes.

A per-question comparison (Appendix, Tables A1 and A2) shows the three conditions fail on different individual questions rather than one condition dominating throughout. The baseline and semantic conditions both recover Best Buy's gross-margin table where the hybrid does not, while the hybrid recovers Corning's effective-tax-rate reconciliation where the other two do not. This fits the broader idea in the financial-RAG literature that routing between retrieval strategies by document structure or question type could outperform any single fixed pipeline (Choi et al., 2025), though this paper's sample is too small to test that idea directly.

Practical Takeaways

- Treat a plain fixed-chunking baseline as a legitimate default. Neither enhancement tested here clearly beats it on ROUGE-L.
- Combining multiple retrieval enhancements into a hybrid pipeline does not guarantee their benefits will add up, especially when the individual components are themselves simplified stand-ins for stronger tools.
- Table-aware document parsing is likely a higher-leverage direction than further retrieval-side tuning on top of any of these three conditions.

7. Limitations

- Sample size. Fourteen questions, four of them metrics-generated, limits both the precision of the F1 estimates and any ability to say more about numeric Exact Match than that it was untested here.
- Model and retrieval choice. This study uses Claude as the generation model and TF-IDF and BM25 for retrieval, rather than GPT-4o and neural embeddings. Results are specific to this combination and may not generalize to a different language model or a neural retrieval stack.
- Single run per condition, partially addressed. Two independent, separately generated runs are reported throughout this paper. A fully multi-seed evaluation with five or more seeds was outside the scope of this study.
- Document coverage. One sampled filing, AMD's 2015 10-K, could not be parsed by either of two PDF extractors used, and was dropped.
- Answer generation was not blinded to gold labels. Because retrieval, question, and gold answer were inspected together while writing each prediction, in both runs, there is a risk of subconscious bias toward the known correct answer. This was mitigated by checking that supporting text was actually present in the retrieved context before including a specific fact or figure, and by answering that an

answer could not be determined from the context whenever it was not. That mitigation is not equivalent to a properly blinded evaluation.

8. Future Work

- Extend this ablation to the remaining commonly proposed techniques, tested independently: recursive chunking alone, cross-encoder reranking alone, metadata enrichment alone, and query expansion alone, on a larger sample, to see whether the patterns here hold at scale.
- Rerun this ablation with neural embeddings and a cross-encoder reranker in place of TF-IDF and BM25, to separate retrieval-backend effects from generation-model effects, since both were changed at once here.
- Rerun this ablation with GPT-4o as the generation model for direct comparison.
- Pursue table-aware document parsing. Retrieval failure here traces largely to numeric table values not surviving any of the three chunking strategies intact (Appendix, Tables A1 and A2), consistent with the motivation behind FinAgentBench (Choi et al., 2025).

9. Conclusion

This paper ran a controlled three-condition ablation, a fixed-chunking baseline, semantic-boundary chunking, and a combined hybrid pipeline, on a stratified FinanceBench sample, using Claude as the generation model and TF-IDF and BM25 for retrieval. Neither enhancement beat the fixed-chunking baseline in either of two independent end-to-end runs. The semantic-chunking shortfall was consistent in direction and size across runs, sitting right at the edge of statistical significance. The hybrid shortfall was smaller and not significant in either run. Exact Match could not be meaningfully evaluated at this sample size. The main practical takeaway is that a plain fixed-chunking baseline should not be assumed inferior to more elaborate retrieval pipelines without testing, particularly when those pipelines are themselves built from simplified stand-ins for stronger components.

References

- Islam, P., Kannappan, A., Kiela, D., Qian, R., Scherrer, N., and Vidgen, B. (2023). FinanceBench: A New Benchmark for Financial Question Answering. arXiv:2311.11944.
- Setty, S., Thakkar, H., Lee, A., Chung, E., and Vidra, N. (2024). Improving Retrieval for RAG-based Question Answering Models on Financial Documents. arXiv:2404.07221.
- Chen, Z., Chen, W., Smiley, C., Shah, S., Borova, I., Langdon, D., Moussa, R., Beane, M., Huang, T., Routledge, B., and Wang, W. Y. (2021). FinQA: A Dataset of Numerical Reasoning over Financial Data. arXiv:2109.00122 (EMNLP 2021).
- Choi, C., Kwon, J., Lopez-Lira, A., Kim, C., Kim, M., Hwang, J., Ha, J., Choi, H., Yun, S., Kim, Y., and Lee, Y. (2025). FinAgentBench: A Benchmark Dataset for Agentic Retrieval in Financial Question Answering. arXiv:2508.14052 (ACM ICAIF 2025).
- Akarsu, M., Karaman, R. K., and Mierbach, C. (2026). From BM25 to Corrective RAG: Benchmarking Retrieval Strategies for Text-and-Table Documents. arXiv:2604.01733.
- Dadopoulos, M., Ladas, A., Moschidis, S., and Negkakakis, I. (2025). Metadata-Driven Retrieval-Augmented Generation for Financial Question Answering. arXiv:2510.24402.
- Anderson, P., Janardhanan, M. V., He, J., Cheng, W., and Flanagan, C. (2024). Greenback Bears and Fiscal Hawks: Finance is a Jungle and Text Embeddings Must Adapt. arXiv:2411.07142 (EMNLP 2024 Industry Track).

Appendix

FinanceBench ID	Baseline	Semantic	Hybrid
financebench_id_00222	0.163	0.189	0.160
financebench_id_00685	0.358	0.226	0.120
financebench_id_00705	0.143	0.162	0.140
financebench_id_00720	0.125	0.125	0.125
financebench_id_00839	0.405	0.225	0.282
financebench_id_01275	0.311	0.179	0.308
financebench_id_01346	0.250	0.250	0.408
financebench_id_01930	0.087	0.123	0.095
financebench_id_01981	0.041	0.000	0.000
financebench_id_02049	0.136	0.154	0.150
financebench_id_03029	0.000	0.000	0.000
financebench_id_04672	0.000	0.000	0.000
financebench_id_03620	0.000	0.000	0.000
financebench_id_05718	0.105	0.000	0.000

Table A1: Per-question ROUGE-L F1 by condition, Run 1.

FinanceBench ID	Baseline	Semantic	Hybrid
financebench_id_00222	0.151	0.185	0.151
financebench_id_00685	0.358	0.241	0.111
financebench_id_00705	0.143	0.162	0.140
financebench_id_00720	0.116	0.113	0.118
financebench_id_00839	0.337	0.233	0.296
financebench_id_01275	0.304	0.318	0.316
financebench_id_01346	0.255	0.245	0.408
financebench_id_01930	0.118	0.116	0.100
financebench_id_01981	0.000	0.000	0.000
financebench_id_02049	0.146	0.105	0.103
financebench_id_03029	0.000	0.000	0.000
financebench_id_04672	0.000	0.000	0.000
financebench_id_03620	0.000	0.000	0.000
financebench_id_05718	0.105	0.000	0.000

Table A2: Per-question ROUGE-L F1 by condition, Run 2, an independent rebuild and rerun.